

Homework 3: Contrastive Divergence and Noise Contrastive Estimation

Due Date: **June 16, 2022**

Submission Guidelines

- Submission only in **pairs** on the course website (Moodle).
- Working environment:
 - We encourage you to work in Jupyter Notebook online using **Google Colab** as it does not require any installation.
- You should submit two **separated** files:
 - A Jupyter file, with the name: ee048954_hw3_id1_id2_id3.py which contains your code implementations.
 - A pdf file, with the name: ee048954_hw3_id1_id2_id3.pdf, which is your report containing plots, answers, and discussions.
- **No handwritten submissions** and no other file-types (.docx, .html...) will be accepted.

Mounting your drive for saving/loading stuff

```
from google.colab import drive
drive.mount('/content/drive')
```

Importing relevant libraries

```
## Standard libraries
import os
import math
import time
import numpy as np
import random
import copy

# Scipy optimization routines
from scipy.optimize import minimize

import tqdm

# Imports for plotting
import matplotlib.pyplot as plt
import matplotlib.animation as animation
from matplotlib import cm
import matplotlib
matplotlib.rcParams['lines.linewidth'] = 2.0
plt.style.use('ggplot')
```

Part I: Contrastive Divergence (50 points)

Consider the following Gaussian Mixture Model (GMM) distributions

$$p(x; \{\mu_i\}) = \sum_{i=1}^K \frac{1}{N} \frac{1}{2\pi} \exp\left\{-\frac{1}{2}\|x - \mu_i\|^2\right\},$$

where $x, \mu_i \in \mathbb{R}^2$. We will use $N = 4$, $\sigma = 1$ and $\{\mu_i\} = \{(0,0)^T, (0,3)^T, (3,0)^T, (3,3)^T\}$.

Sampling from GMM

Task 1 Direct sampling: Use your function from HW1 that accepts $\{\mu_i\}$, and returns a sample x from $p(x; \{\mu_i\})$. Draw $J = 1000$ samples $\{x_j\}$ from the distribution $p(x; \{\mu_i\})$ using this function. These will be our **real samples**.

```
def plot_scatter(samples, title, group=None):
    if group is not None:
        plt.scatter(samples[:, 0], samples[:, 1], c=group)
    else:
        plt.scatter(samples[:, 0], samples[:, 1])
    plt.title(title)
    plt.show()
```

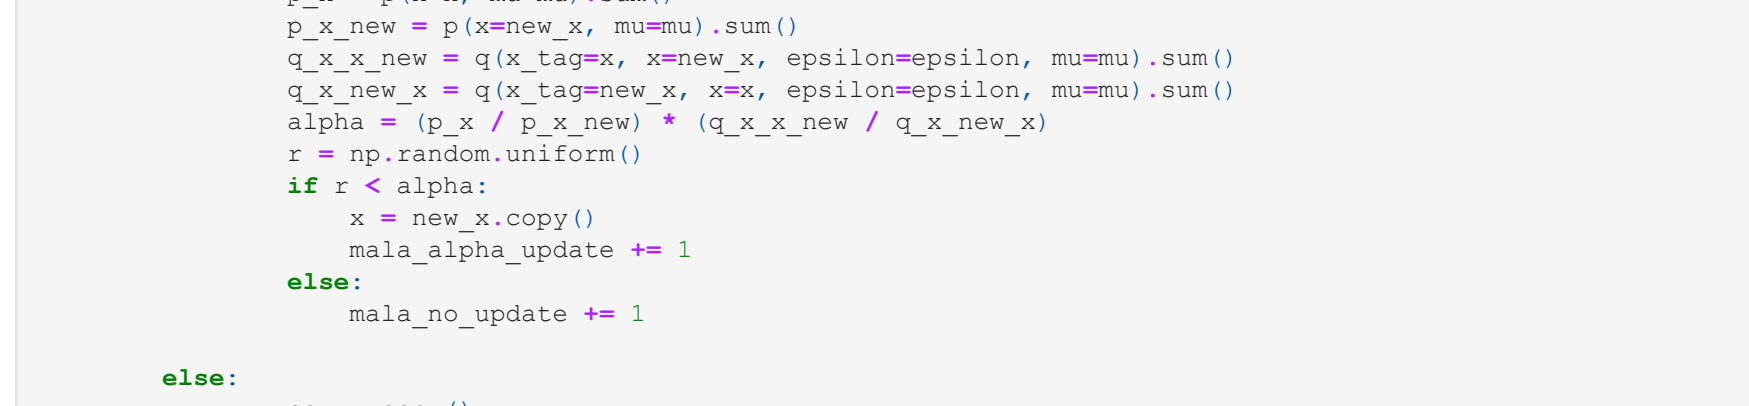
```
def mix_plot_scatter(samples_list, titles_list, width, figsize=(15,5)):
    N = len(samples_list)
    fig, axes = plt.subplots(1, 2, figsize=figsize)
    rows = int(N/width)
    for i in range(rows):
        for j in range(width):
            idx = rows * i + j
            if rows > 1:
                x = np.random.choice(samples_list[idx][:, 0], samples_list[idx][:, 1])
                axes[j,1].set_title(titles_list[idx])
            else:
                axes[idx].set_title(titles_list[idx])
                axes[idx].scatter(samples_list[idx][:, 0], samples_list[idx][:, 1])
```

```
def mix_gauss_draw(N=4, sigma=1, mu=[np.array([0,0]),
                                         np.array([0,3]),
                                         np.array([3,0]),
                                         np.array([3,3])], J=1_000):
    Sigma = np.eye(len(mu[0])) * (sigma ** 2)
    X = np.zeros((J, len(mu[0])))
    G = np.zeros((J, 1))

    for i in range(J):
        mu = np.random.choice(N, 1)
        M = mu[int(mu)]
        x = np.random.multivariate_normal(mean=M, cov=Sigma)
        X[i,:] = x
        G[i] = mu

    return X, G
```

```
real_samples, groups = mix_gauss_draw(sigma=1)
```



Task 2 Sampling with MCMC: Implement the MALA algorithm to draw samples from $p(x; \{\mu_i\})$. The function will get an initial guess $\{x_i\}$ and will generate chains of length L . Use $\sqrt{2} = 0.1$ and $N \sim \mathcal{N}(0, I)$.

```
def log_gradient_p(x, sigma=2, mu=[np.array([0,0]), np.array([0,3]),
                                         np.array([3,0]), np.array([3,3])]):
    N = len(mu)
    denominator = 0
    for i in range(N):
        denominator += np.exp(-(1/(2*sigma**2)) * np.linalg.norm(x-mu[i], axis=1) ** 2)
    denominator = np.reshape(denominator, (denominator.shape[0],1))

    nominator = 0
    for i in range(N):
        s = np.exp(-(1/(2*sigma**2)) * np.linalg.norm(x-mu[i], axis=1) ** 2)
        q_x_new_x = q(x-tagnew_x, sigma, epsilon=epsilon, mu=mu)
        alpha = (p(x) / p_x_new) * (q_x_new_x / q_x_new_x)
        nominator += s * d

    return (1/sigma**2) * nominator.transpose() / denominator
```

```
def p(x, sigma=1, mu=[np.array([0,0]), np.array([0,3]),
                                         np.array([3,0]), np.array([3,3])]):
    N = len(mu)
    sum_exp = 0
    for n in range(N):
        log_exp = -0.5 * (np.linalg.norm(x-mu[n], axis=1) ** 2)
        mu_exp = normf((sigma))
        return (1/N) * (1 / (2*np.pi)) * sum_exp
```

```
def q(x_tag, x, epsilon, mu=[np.array([0,0]), np.array([3,3]),
                                         np.array([0,3]), np.array([3,0])]):
    norm = np.linalg.norm(x_tag - x - epsilon * log_gradient_p(x, mu=mu), axis=1) ** 2
    log_q = normf((epsilon))
    q = np.exp(log_q)
    return q
```

```
def langevin_mala_dynamics(initial_guess, steps=5000, samples=1000, sigma=2, factored_epsilon=0.1,
                             mu=[np.array([0,0]), np.array([0,3]), np.array([3,0]), np.array([3,3])],
                             MALA=False, verbose=False):
    epsilon = (factored_epsilon ** 2) / 2
    x = initial_guess
    mala_reg_updates = 0
    mala_alpha_update = 0
    mala_no_update = 0
    for i in tqdm.tqdm(range(steps)):
        noise = np.random.randn(x.shape[0], x.shape[1])
        new_x = x + epsilon * log_gradient_p(x, sigma=2*sigma, mu=mu) + factored_epsilon * noise

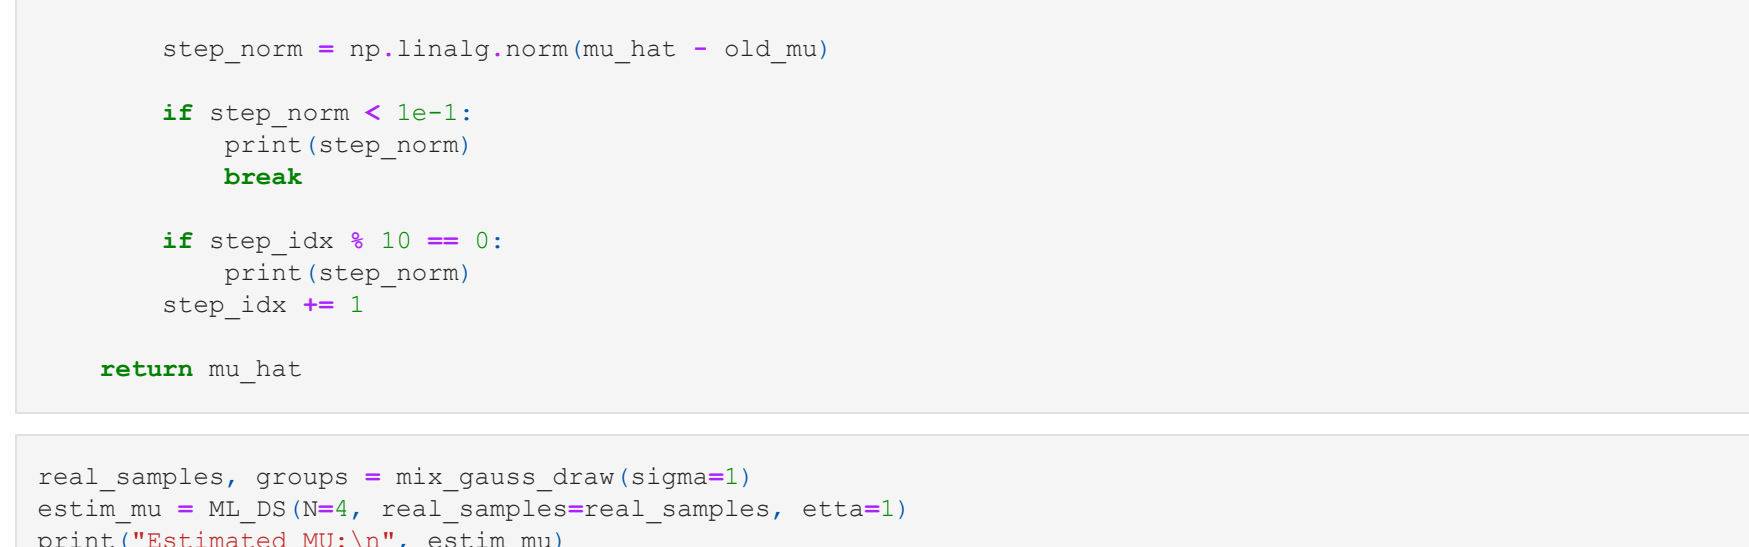
        if MALA:
            if p(new_x, mu=mu).sum() > p(x, mu=mu).sum():
                x = new_x.copy()
                mala_reg_updates += 1
            else:
                x = p(x, mu=mu).sum()
                p_x_new = p(x, mu=mu).sum()
                q_x_new_x = q(x_tagnew_x, sigma, epsilon=epsilon, mu=mu).sum()
                alpha = (p(x) / p_x_new) * (q_x_new_x / q_x_new_x)
                if alpha > 1:
                    x = new_x.copy()
                    mala_alpha_update += 1
                else:
                    mala_no_update += 1
                    x = new_x.copy()

        if MALA and verbose:
            print(mala_reg_updates / steps)
            print(mala_alpha_update / steps)
            print(mala_no_update / steps)

    return x
```

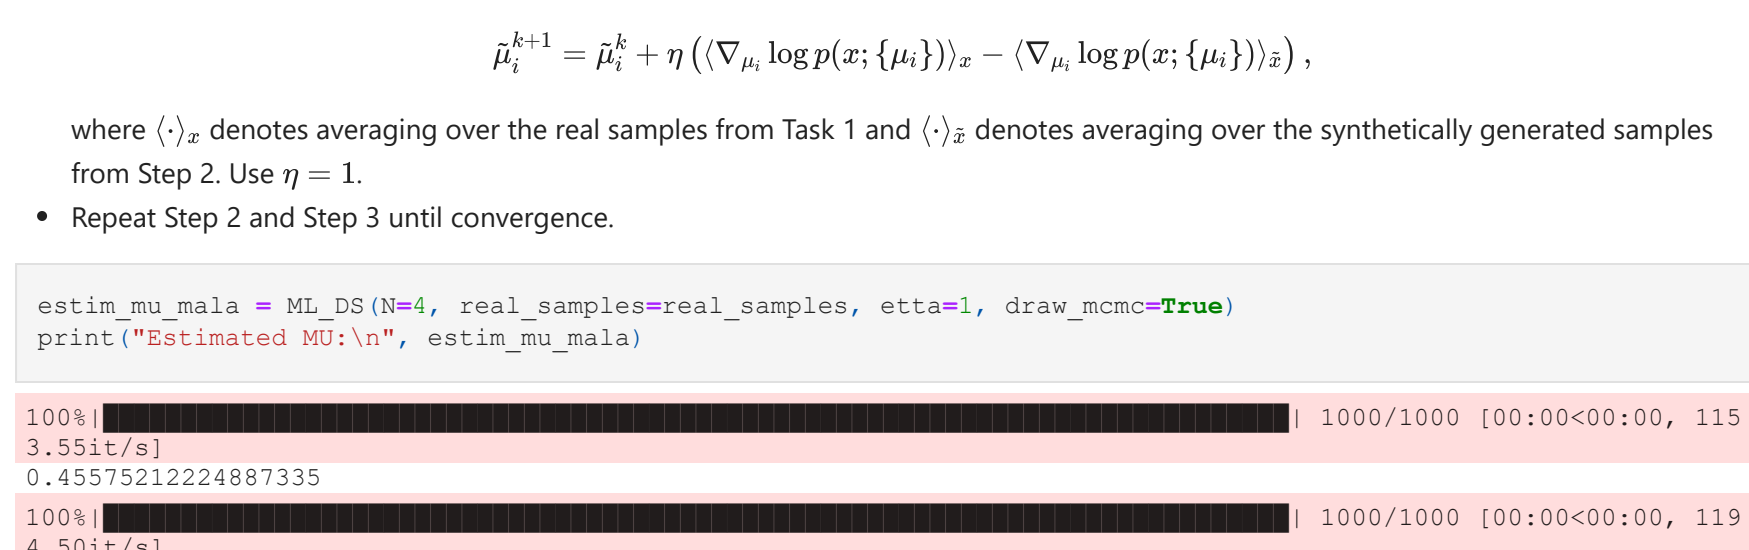
```
initial_guess = np.random.uniform(low=0, high=3, size=(1_000, 2))
x = langevin_mala_dynamics(initial_guess=initial_guess, steps=10_000, samples=1_000, MALA=False, sigma=2)

100% | 10000/10000 | 00:05:00:00, 180 9.22it/s
```



```
initial_guess = np.random.uniform(low=0, high=3, size=(1_000, 2))
x = langevin_mala_dynamics(initial_guess=initial_guess, steps=10_000, samples=1_000, MALA=True, sigma=2)

100% | 10000/10000 | 00:05:00:00, 60 9.22it/s
```



From now on, we will refer to $\{\mu_i\}$ as **unknowns** and we will estimate them using different algorithms.

Estimation of $\{\mu_i\}$

Task 3 Implement Maximum likelihood (ML) estimation of $\{\mu_i\}$ using direct sampling:

- Step 1: Randomly initialize $\{\hat{\mu}_i\}$ from $U([0,3]^2)$.
- Step 2: Use your function from Task 1 to draw 100 samples \hat{x} from $p(x; \{\mu_i\})$ using $\{\hat{\mu}_i\}$.
- Step 3: Update $\{\hat{\mu}_i\}$ using the ML gradient descent step:

$$\hat{\mu}_i^{t+1} = \hat{\mu}_i^t + \eta \left((\nabla_{\mu_i} \log p(x; \{\mu_i\}))_x - (\nabla_{\mu_i} \log p(x; \{\mu_i\}))_z \right),$$

where $(\cdot)_x$ denotes averaging over the real samples from Task 1 and $(\cdot)_z$ denotes averaging over the synthetically generated samples from Step 2. Use $\eta = 1$.

- Repeat Step 2 and Step 3 until convergence.

```
def grad_mu_log_p(x, mu):
    N = len(mu)
    denominator = 0
    for n in range(N):
        log_exp = -0.5 * (np.linalg.norm(x-mu[n], axis=1) ** 2)
        denominator += np.exp(log_exp)

    nominator = np.zeros((N,2))
    for n in range(N):
        log_exp = -0.5 * (np.linalg.norm(x-mu[n], axis=1) ** 2)
        exp = np.exp(log_exp)
        nominator[n] = exp * (x-mu[n])

    return (1/denominator) * nominator
```

```
def avg_grad_mu_log_p(x, mu):
    grad = grad_mu_log_p(x, mu)
    return grad / x.shape[0]
```

```
def ML_DS(N, real_samples, etta=1, draw_mcmc=False, initial_task=1):
    mu_hat = np.random.uniform(size=(N,2)) * 3
    step_idx = 0
    while True:
        if draw_mcmc:
            if initial_task_1:
                initial_guess = mix_gauss_draw(N=N, sigma=1, J=100) # using original mu's
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=10, samples=100, MALA=True, sigma=2)
            else:
                initial_guess = np.random.randn(100, 2) * np.sqrt(2) * 1.5
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=1_000, samples=100, MALA=True, sigma=2)
            x_hat = mix_gauss_draw(N=N, sigma=1, mu=mu_hat, J=100)
            avg_mu_grad_estim = avg_grad_mu_log_p(x_hat, mu_hat)
            avg_mu_grad_real = avg_grad_mu_log_p(real_samples, mu_hat)

            old_mu = mu_hat.copy()
            mu_hat = mu_hat + etta * (avg_mu_grad_real - avg_mu_grad_estim)

            step_norm = np.linalg.norm(mu_hat - old_mu)

            if step_norm < 1e-1:
                print('step_norm')
                break

            if step_idx % 10 == 0:
                print('step_norm')
                step_idx += 1

    return mu_hat
```

```
real_samples, groups = mix_gauss_draw(sigma=1)
estim_mu = ML_DS(N=4, real_samples=real_samples, etta=1)
print("Estimated MU:\n", estim_mu)

0.2558766759562691
0.2110447927951567
[-0.38759686 -0.44497931]
0.0880344255979092
[[-0.31091859 0.03061885]
 [-0.01691787 3.00878443]
 [-0.981965 2.99639136]
 [-0.0769686 -0.04974094]]
```

Task 4 Implement Maximum likelihood (ML) estimation of $\{\mu_i\}$ using MCMC:

- Step 1: Randomly initialize $\{\hat{\mu}_i\}$ from $U([0,3]^2)$.
- Step 2: Use your function from Task 2 to draw 100 samples \hat{x} from $p(x; \{\mu_i\})$ using $\{\hat{\mu}_i\}$. Initialize the chains with $\hat{x}_i \sim \mathcal{N}(1.5, 2)$ and use chains length of $L = 1000$.
- Step 3: Update $\{\hat{\mu}_i\}$ using the ML gradient descent step:

$$\hat{\mu}_i^{t+1} = \hat{\mu}_i^t + \eta \left((\nabla_{\mu_i} \log p(x; \{\mu_i\}))_x - (\nabla_{\mu_i} \log p(x; \{\mu_i\}))_z \right),$$

where $(\cdot)_x$ denotes averaging over the real samples from Task 1 and $(\cdot)_z$ denotes averaging over the synthetically generated samples from Step 2. Use $\eta = 1$.

- Repeat Step 2 and Step 3 until convergence.

```
estim_mu_mala = ML_DS(N=4, real_samples=real_samples, etta=1, draw_mcmc=True)
print("Estimated MU:\n", estim_mu_mala)

100% | 10000/10000 | 00:00:00:00, 113 6.84it/s
0.2558766759562691
0.2110447927951567
[-0.38759686 -0.44497931]
0.0880344255979092
[[-0.31091859 0.03061885]
 [-0.01691787 3.00878443]
 [-0.981965 2.99639136]
 [-0.0769686 -0.04974094]]
```

```
def CD_estim_mu_log_p(x, mu):
    N = len(mu)
    denominator = 0
    for n in range(N):
        log_exp = -0.5 * (np.linalg.norm(x-mu[n], axis=1) ** 2)
        denominator += np.exp(log_exp)

    nominator = np.zeros((N,2))
    for n in range(N):
        log_exp = -0.5 * (np.linalg.norm(x-mu[n], axis=1) ** 2)
        exp = np.exp(log_exp)
        nominator[n] = exp * (x-mu[n])

    return (1/denominator) * nominator
```

```
def avg_grad_mu_log_p(x, mu):
    grad = grad_mu_log_p(x, mu)
    return grad / x.shape[0]
```

```
def CD_DS(N, real_samples, etta=1, draw_mcmc=True, initial_task=1):
    mu_hat = np.random.uniform(size=(N,2)) * 3
    step_idx = 0
    while True:
        if draw_mcmc:
            if initial_task_1:
                initial_guess = mix_gauss_draw(N=N, sigma=1, J=100) # using original mu's
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=10, samples=100, MALA=True, sigma=2)
            else:
                initial_guess = np.random.randn(100, 2) * np.sqrt(2) * 1.5
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=1_000, samples=100, MALA=True, sigma=2)
            x_hat = mix_gauss_draw(N=N, sigma=1, mu=mu_hat, J=100)
            avg_mu_grad_estim = avg_grad_mu_log_p(x_hat, mu_hat)
            avg_mu_grad_real = avg_grad_mu_log_p(real_samples, mu_hat)

            old_mu = mu_hat.copy()
            mu_hat = mu_hat + etta * (avg_mu_grad_real - avg_mu_grad_estim)

            step_norm = np.linalg.norm(mu_hat - old_mu)

            if step_norm < 1e-1:
                print('step_norm')
                break

            if step_idx % 10 == 0:
                print('step_norm')
                step_idx += 1

    return mu_hat
```

```
real_samples, groups = mix_gauss_draw(sigma=1)
estim_mu = CD_DS(N=4, real_samples=real_samples, etta=1)
print("Estimated MU:\n", estim_mu)

0.2558766759562691
0.2110447927951567
[-0.38759686 -0.44497931]
0.0880344255979092
[[-0.31091859 0.03061885]
 [-0.01691787 3.00878443]
 [-0.981965 2.99639136]
 [-0.0769686 -0.04974094]]
```

```
def CD_DS(N, real_samples, etta=1, draw_mcmc=True, initial_task=1):
    mu_hat = np.random.uniform(size=(N,2)) * 3
    step_idx = 0
    while True:
        if draw_mcmc:
            if initial_task_1:
                initial_guess = mix_gauss_draw(N=N, sigma=1, J=100) # using original mu's
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=10, samples=100, MALA=True, sigma=2)
            else:
                initial_guess = np.random.randn(100, 2) * np.sqrt(2) * 1.5
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=1_000, samples=100, MALA=True, sigma=2)
            x_hat = mix_gauss_draw(N=N, sigma=1, mu=mu_hat, J=100)
            avg_mu_grad_estim = avg_grad_mu_log_p(x_hat, mu_hat)
            avg_mu_grad_real = avg_grad_mu_log_p(real_samples, mu_hat)

            old_mu = mu_hat.copy()
            mu_hat = mu_hat + etta * (avg_mu_grad_real - avg_mu_grad_estim)

            step_norm = np.linalg.norm(mu_hat - old_mu)

            if step_norm < 1e-1:
                print('step_norm')
                break

            if step_idx % 10 == 0:
                print('step_norm')
                step_idx += 1

    return mu_hat
```

```
real_samples, groups = mix_gauss_draw(sigma=1)
estim_mu = CD_DS(N=4, real_samples=real_samples, etta=1)
print("Estimated MU:\n", estim_mu)

0.2558766759562691
0.2110447927951567
[-0.38759686 -0.44497931]
0.0880344255979092
[[-0.31091859 0.03061885]
 [-0.01691787 3.00878443]
 [-0.981965 2.99639136]
 [-0.0769686 -0.04974094]]
```

```
def CD_DS(N, real_samples, etta=1, draw_mcmc=True, initial_task=1):
    mu_hat = np.random.uniform(size=(N,2)) * 3
    step_idx = 0
    while True:
        if draw_mcmc:
            if initial_task_1:
                initial_guess = mix_gauss_draw(N=N, sigma=1, J=100) # using original mu's
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=10, samples=100, MALA=True, sigma=2)
            else:
                initial_guess = np.random.randn(100, 2) * np.sqrt(2) * 1.5
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=1_000, samples=100, MALA=True, sigma=2)
            x_hat = mix_gauss_draw(N=N, sigma=1, mu=mu_hat, J=100)
            avg_mu_grad_estim = avg_grad_mu_log_p(x_hat, mu_hat)
            avg_mu_grad_real = avg_grad_mu_log_p(real_samples, mu_hat)

            old_mu = mu_hat.copy()
            mu_hat = mu_hat + etta * (avg_mu_grad_real - avg_mu_grad_estim)

            step_norm = np.linalg.norm(mu_hat - old_mu)

            if step_norm < 1e-1:
                print('step_norm')
                break

            if step_idx % 10 == 0:
                print('step_norm')
                step_idx += 1

    return mu_hat
```

```
real_samples, groups = mix_gauss_draw(sigma=1)
estim_mu = CD_DS(N=4, real_samples=real_samples, etta=1)
print("Estimated MU:\n", estim_mu)

0.2558766759562691
0.2110447927951567
[-0.38759686 -0.44497931]
0.0880344255979092
[[-0.31091859 0.03061885]
 [-0.01691787 3.00878443]
 [-0.981965 2.99639136]
 [-0.0769686 -0.04974094]]
```

```
def CD_DS(N, real_samples, etta=1, draw_mcmc=True, initial_task=1):
    mu_hat = np.random.uniform(size=(N,2)) * 3
    step_idx = 0
    while True:
        if draw_mcmc:
            if initial_task_1:
                initial_guess = mix_gauss_draw(N=N, sigma=1, J=100) # using original mu's
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=10, samples=100, MALA=True, sigma=2)
            else:
                initial_guess = np.random.randn(100, 2) * np.sqrt(2) * 1.5
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=1_000, samples=100, MALA=True, sigma=2)
            x_hat = mix_gauss_draw(N=N, sigma=1, mu=mu_hat, J=100)
            avg_mu_grad_estim = avg_grad_mu_log_p(x_hat, mu_hat)
            avg_mu_grad_real = avg_grad_mu_log_p(real_samples, mu_hat)

            old_mu = mu_hat.copy()
            mu_hat = mu_hat + etta * (avg_mu_grad_real - avg_mu_grad_estim)

            step_norm = np.linalg.norm(mu_hat - old_mu)

            if step_norm < 1e-1:
                print('step_norm')
                break

            if step_idx % 10 == 0:
                print('step_norm')
                step_idx += 1

    return mu_hat
```

```
real_samples, groups = mix_gauss_draw(sigma=1)
estim_mu = CD_DS(N=4, real_samples=real_samples, etta=1)
print("Estimated MU:\n", estim_mu)

0.2558766759562691
0.2110447927951567
[-0.38759686 -0.44497931]
0.0880344255979092
[[-0.31091859 0.03061885]
 [-0.01691787 3.00878443]
 [-0.981965 2.99639136]
 [-0.0769686 -0.04974094]]
```

```
def CD_DS(N, real_samples, etta=1, draw_mcmc=True, initial_task=1):
    mu_hat = np.random.uniform(size=(N,2)) * 3
    step_idx = 0
    while True:
        if draw_mcmc:
            if initial_task_1:
                initial_guess = mix_gauss_draw(N=N, sigma=1, J=100) # using original mu's
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=10, samples=100, MALA=True, sigma=2)
            else:
                initial_guess = np.random.randn(100, 2) * np.sqrt(2) * 1.5
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=1_000, samples=100, MALA=True, sigma=2)
            x_hat = mix_gauss_draw(N=N, sigma=1, mu=mu_hat, J=100)
            avg_mu_grad_estim = avg_grad_mu_log_p(x_hat, mu_hat)
            avg_mu_grad_real = avg_grad_mu_log_p(real_samples, mu_hat)

            old_mu = mu_hat.copy()
            mu_hat = mu_hat + etta * (avg_mu_grad_real - avg_mu_grad_estim)

            step_norm = np.linalg.norm(mu_hat - old_mu)

            if step_norm < 1e-1:
                print('step_norm')
                break

            if step_idx % 10 == 0:
                print('step_norm')
                step_idx += 1

    return mu_hat
```

```
real_samples, groups = mix_gauss_draw(sigma=1)
estim_mu = CD_DS(N=4, real_samples=real_samples, etta=1)
print("Estimated MU:\n", estim_mu)

0.2558766759562691
0.2110447927951567
[-0.38759686 -0.44497931]
0.0880344255979092
[[-0.31091859 0.03061885]
 [-0.01691787 3.00878443]
 [-0.981965 2.99639136]
 [-0.0769686 -0.04974094]]
```

```
def CD_DS(N, real_samples, etta=1, draw_mcmc=True, initial_task=1):
    mu_hat = np.random.uniform(size=(N,2)) * 3
    step_idx = 0
    while True:
        if draw_mcmc:
            if initial_task_1:
                initial_guess = mix_gauss_draw(N=N, sigma=1, J=100) # using original mu's
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=10, samples=100, MALA=True, sigma=2)
            else:
                initial_guess = np.random.randn(100, 2) * np.sqrt(2) * 1.5
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=1_000, samples=100, MALA=True, sigma=2)
            x_hat = mix_gauss_draw(N=N, sigma=1, mu=mu_hat, J=100)
            avg_mu_grad_estim = avg_grad_mu_log_p(x_hat, mu_hat)
            avg_mu_grad_real = avg_grad_mu_log_p(real_samples, mu_hat)

            old_mu = mu_hat.copy()
            mu_hat = mu_hat + etta * (avg_mu_grad_real - avg_mu_grad_estim)

            step_norm = np.linalg.norm(mu_hat - old_mu)

            if step_norm < 1e-1:
                print('step_norm')
                break

            if step_idx % 10 == 0:
                print('step_norm')
                step_idx += 1

    return mu_hat
```

```
real_samples, groups = mix_gauss_draw(sigma=1)
estim_mu = CD_DS(N=4, real_samples=real_samples, etta=1)
print("Estimated MU:\n", estim_mu)

0.2558766759562691
0.2110447927951567
[-0.38759686 -0.44497931]
0.0880344255979092
[[-0.31091859 0.03061885]
 [-0.01691787 3.00878443]
 [-0.981965 2.99639136]
 [-0.0769686 -0.04974094]]
```

```
def CD_DS(N, real_samples, etta=1, draw_mcmc=True, initial_task=1):
    mu_hat = np.random.uniform(size=(N,2)) * 3
    step_idx = 0
    while True:
        if draw_mcmc:
            if initial_task_1:
                initial_guess = mix_gauss_draw(N=N, sigma=1, J=100) # using original mu's
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=10, samples=100, MALA=True, sigma=2)
            else:
                initial_guess = np.random.randn(100, 2) * np.sqrt(2) * 1.5
                mu_hat = langevin_mala_dynamics(initial_guess=initial_guess, steps=1_000, samples=100, MALA=True, sigma=2)
            x_hat = mix_gauss_draw(N=N, sigma=1, mu=mu_hat, J=100)
            avg_mu_grad_estim = avg_grad_mu_log_p(x_hat, mu_hat)
            avg_mu_grad_real = avg_grad_mu_log_p(real_samples, mu_hat)

            old_mu = mu_hat.copy()
            mu_hat = mu_hat + etta * (avg_mu_grad_real - avg_mu_grad_estim)

            step_norm = np.linalg.norm(mu_hat - old_mu)

            if step_norm < 1e-1:
                print('step_norm')
                break

            if step_idx % 10 == 0:
                print('step_norm')
                step_idx += 1

    return mu_hat
```

```
real_samples, groups = mix_gauss_draw(sigma=1)
estim_mu = CD_DS(N=4, real_samples=real_samples, etta=1)
print("Estimated MU:\n", estim_mu)

0.2558766759562691
0.2110447927951567
[-0.38759686 -0.44497931]
0.0880344255979092
[[-0.31091859 0.03061885]
 [-0.01691787 3.00878443]
 [-0.981965 2.99639136]
 [-0.0769686 -0.04974094]]
```




```
100% | 5/50 [00:01<00:08, 5.01it/s]C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: divide by zero encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\scipy\optimize\optimize.py:663: RuntimeWarning: invalid value encountered in double_scalars
grad[k] = f(*(xx + d,) + args)) - f0) / d[k]
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: divide by zero encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: invalid value encountered in log
56% | 28/50 [00:05<00:05, 4.15it/s]C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: invalid value encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: divide by zero encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\scipy\optimize\optimize.py:663: RuntimeWarning: invalid value encountered in double_scalars
grad[k] = f(*(xx + d,) + args)) - f0) / d[k]
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: invalid value encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: divide by zero encountered in log
100% | 50/50 [00:10<00:00, 4.65it/s]
44% | 22/50 [00:04<00:06, 4.65it/s]C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: invalid value encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: divide by zero encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\scipy\optimize\optimize.py:663: RuntimeWarning: invalid value encountered in double_scalars
grad[k] = f(*(xx + d,) + args)) - f0) / d[k]
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: invalid value encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: divide by zero encountered in log
100% | 50/50 [00:11<00:00, 4.35it/s]
22% | 11/50 [00:07<00:25, 1.55it/s]C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: divide by zero encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\scipy\optimize\optimize.py:663: RuntimeWarning: invalid value encountered in double_scalars
grad[k] = f(*(xx + d,) + args)) - f0) / d[k]
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: divide by zero encountered in log
C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\ipykernel_launcher.py:6: RuntimeWarning: invalid value encountered in log
100% | 50/50 [00:34<00:00, 1.47it/s]
100% | 50/50 [01:01<00:00, 1.23it/s]
100% | 50/50 [01:26<00:00, 1.74it/s]
74% | 37/50 [01:40<00:38, 2.96s/it]C:\Users\neria\anaconda3\envs\deep_learn\lib\site-packages\scipy\optimize\optimize.py:1013: RuntimeWarning: divide by zero encountered in double_scalars
rhok = 1.0 / (numpy.dot(yk, sk))
100% | 50/50 [02:16<00:00, 2.74s/it]
```

```
In [105]: nine_task_j_acc = {}
ten_task_j_acc = {}
for key, val in dic_res.items():
    print("task:", key[0], "J=", key[1])
    print("Failures:", val[0], "Avg Err:", val[1]/(50-val[0]))

    if key[0] == 9:
        nine_task_j_acc[key[1]] = val[1]/(50-val[0])
    if key[0] == 10:
        ten_task_j_acc[key[1]] = val[1]/(50-val[0])

    print("====")

Task: 9 J= 100
Failures: 2 Avg Err: 2.5393292646359957
=====
Task: 10 J= 100
Failures: 20 Avg Err: 566.2526486742096
=====
Task: 9 J= 150
Failures: 0 Avg Err: 1.6880837649611724
=====
Task: 10 J= 150
Failures: 26 Avg Err: 2646.624648537139
=====
Task: 9 J= 1000
Failures: 0 Avg Err: 0.7642552780696296
=====
Task: 10 J= 1000
Failures: 24 Avg Err: 14442.679453818524
=====
Task: 9 J= 2000
Failures: 0 Avg Err: 0.696888464449651
=====
Task: 10 J= 2000
Failures: 29 Avg Err: 50842.81310065304
=====
Task: 9 J= 3000
Failures: 0 Avg Err: 0.678179897579023
=====
Task: 10 J= 3000
Failures: 27 Avg Err: 39854.68709802683
=====
Task: 9 J= 5000
Failures: 0 Avg Err: 0.6487374382663149
=====
Task: 10 J= 5000
Failures: 34 Avg Err: 135038.9161081356
=====
```

 **Task 13:** Discussion: How does the number of samples J affect the accuracy of the estimation? How does the addition of Z as an unknown affect the accuracy?

```
In [126]: fig, axes = plt.subplots(1,2, figsize=(15,5))
axes[0].plot(list(ten_task_j_acc.keys()), list(ten_task_j_acc.values()), label="Error")
axes[0].legend()
axes[1].set_title("Task 10")
axes[1].set_xlabel("J")
axes[1].set_ylabel("Avg Err")

axes[0].plot(list(nine_task_j_acc.keys()), list(nine_task_j_acc.values()), label="Error")
axes[0].legend()
axes[0].set_title("Task 9")
axes[0].set_xlabel("J")
axes[0].set_ylabel("Avg Err")

Out[126]: Text(0, 0.5, 'Avg Err')
```



Discussion:

In case we know Z the error decrease with J .

In the other case where Z need to be estimated we both have higher number of errors and worst error due to stability issues.